

# cycle\_gan

December 3, 2025

## 1 Style transfer with Cycle-GAN

I got introduced to the concept through Kaggle's "I'm something of a painter myself" competition. Here's my go at it. The architecture of the generator and the discriminator, and the loss function are inspired by [Amy Jang's notebook](#) and [this paper](#) by Zhu et al.

```
[ ]: import tensorflow as tf
!pip install tensorflow_addons
import tensorflow_addons as tfa
from tensorflow import keras
from tensorflow.keras import layers
import tensorflow_datasets as tfds
import os

import matplotlib.pyplot as plt
import numpy as np
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting tensorflow_addons
  Downloading tensorflow_addons-0.19.0-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.1 MB)
                                             1.1/1.1 MB
14.6 MB/s eta 0:00:00
Requirement already satisfied: typeguard>=2.7 in
/usr/local/lib/python3.8/dist-packages (from tensorflow_addons) (2.7.1)
Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (from tensorflow_addons) (21.3)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in
/usr/local/lib/python3.8/dist-packages (from packaging->tensorflow_addons) (3.0.9)
Installing collected packages: tensorflow_addons
Successfully installed tensorflow_addons-0.19.0
```

### 1.1 Preprocessing the images.

Cast to float32, renormalize to  $[-1, 1]$ , and reshape to  $256 \times 256 \times 3$ .

```
[ ]: IMAGE_SIZE = [256, 256]

def decode_image(image):
    #image = tf.image.decode_jpeg(image, channels=3)
    image = tf.cast(image, tf.float32)
    image = (image / tf.reduce_max(image))*2 - 1 ## Scale the image between -1 and 1 (divide by half the max and then - 1)
    image = tf.reshape(image, [*IMAGE_SIZE, 3])
    return image
```

## 1.2 Download the dataset from tfds

```
[ ]: tfds_data = tfds.builder('cycle_gan', config='monet2photo')
tfds_data.download_and_prepare()
monet_ds = tfds_data.as_dataset(split='trainA', shuffle_files=True,
                                batch_size=-1)[['image']]
photo_ds = tfds_data.as_dataset(split='trainB', shuffle_files=True,
                                batch_size=-1)[['image']]

Downloading and preparing dataset 291.09 MiB (download: 291.09 MiB, generated:
Unknown size, total: 291.09 MiB) to
/root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0...
Dl Completed...: 0 url [00:00, ? url/s]
Dl Size...: 0 MiB [00:00, ? MiB/s]
Extraction completed...: 0 file [00:00, ? file/s]
Generating splits...: 0%|          | 0/4 [00:00<?, ? splits/s]
Generating trainA examples...: 0%|          | 0/1072 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0.incompletePSK8T8/
cycle_gan-trainA.tfrecord*.....
Generating trainB examples...: 0%|          | 0/6287 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0.incompletePSK8T8/
cycle_gan-trainB.tfrecord*.....
Generating testA examples...: 0%|          | 0/121 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0.incompletePSK8T8/
cycle_gan-testA.tfrecord*...:...
Generating testB examples...: 0%|          | 0/751 [00:00<?, ? examples/s]
Shuffling /root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0.incompletePSK8T8/
cycle_gan-testB.tfrecord*...:...
Dataset cycle_gan downloaded and prepared to
/root/tensorflow_datasets/cycle_gan/monet2photo/2.0.0. Subsequent calls will
reuse this data.
```

### 1.3 Setup for TPU usage

I didn't use TPU for training but this should take care of the setup. There might be some changes to the training loop as well since this uses the distributed strategy (`strategy.run`, etc.)

```
[ ]: try:
    tpu = tf.distribute.cluster_resolver.TPUClusterResolver()
    print('Device:', tpu.master())
    tf.config.experimental_connect_to_cluster(tpu)
    tf.tpu.experimental.initialize_tpu_system(tpu)
    strategy = tf.distribute.TPUStrategy(tpu)
except:
    print("TPU setup failed")
    strategy = tf.distribute.get_strategy()
print('Number of replicas:', strategy.num_replicas_in_sync)

AUTOTUNE = tf.data.experimental.AUTOTUNE

print(tf.__version__)
```

```
TPU setup failed
Number of replicas: 1
2.9.2
```

### 1.4 setup the Dataset objects.

This doesn't have to be done exactly like this but the `drop_remainder=True` setting is necessary for XLA and TPU training. Doesn't matter for my case.

```
[ ]: with strategy.scope():
    monet_ds = tf.data.Dataset.from_tensor_slices(monet_ds).map(decode_image).
        batch(1, drop_remainder=True).prefetch(AUTOTUNE)
    photo_ds = tf.data.Dataset.from_tensor_slices(photo_ds).map(decode_image).
        batch(1, drop_remainder=True).prefetch(AUTOTUNE)
```

take one example of each dataset


```
[ ]: example_monet = next(iter(monet_ds))
example_photo = next(iter(photo_ds))
```

have a look at the examples

```
[ ]: fig, ax = plt.subplots(1, 2, dpi=150)
ax[0].set_title('Photo')
ax[0].imshow(example_photo[0]*0.5 + 0.5)

ax[1].set_title('Monet')
ax[1].imshow(example_monet[0]*0.5 + 0.5)
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f53ebf80be0>
```



## 1.5 Downsampling unit

A convolution layer followed by instance normalization and a LeakyReLU activation.

```
[ ]: class Downsample(layers.Layer):
    def __init__(self, filters, apply_InstanceNorm=True, **kwargs):
        super(Downsample, self).__init__(**kwargs)
        self.instance_Norm = apply_InstanceNorm
        self.init = tf.random_normal_initializer(mean=0., stddev=0.02)
        self.init_gamma = keras.initializers.RandomNormal(mean=0., stddev=0.02)
        self.conv = layers.Conv2D(filters, 4, strides=2, padding='same', \
                               use_bias=False, kernel_initializer=self.init)
        self.insNorm = tfa.layers.InstanceNormalization(gamma_initializer=self.
                                                       init_gamma)
        self.activation = layers.LeakyReLU()

    def call(self, inputs):
        x = self.conv(inputs)
        if self.instance_Norm:
            x = self.insNorm(x)
        return self.activation(x)
```

## 1.6 Upsampling unit

A transpose convolution unit followed by instance normalization, dropout, and ReLU activation

```
[ ]: class UpSample(layers.Layer):
    def __init__(self, filters, apply_dropout=False, **kwargs):
        super(UpSample, self).__init__(**kwargs)
        self.apply_dropout = apply_dropout
```

```

    self.init = tf.random_normal_initializer(mean=0., stddev=0.02)
    self.init_gamma = keras.initializers.RandomNormal(mean=0., stddev=0.02)
    self.conv = layers.Conv2DTranspose(filters, 4, strides=2, ↴
padding='same', \
                                     use_bias=False, kernel_initializer=self.init)
    self.insNorm = tfa.layers.InstanceNormalization(gamma_initializer=self. ↴
init_gamma)
    self.dropout = layers.Dropout(0.5)
    self.activation = layers.ReLU()

    def call(self, inputs):
        x = self.conv(inputs)
        x = self.insNorm(x)
        if self.apply_dropout:
            x = self.dropout(x)
        return self.activation(x)

```

## 1.7 The Generator

The Generator is a U-Net with 8 downsample units (64-128-256-512-512-512-512-512) and 7 upsampling units (512-512-512-512-256-128-64) connected to a final Conv2DTranspose layer.

```
[ ]: class Generator(keras.models.Model):
    def __init__(self, **kwargs):
        super(Generator, self).__init__(**kwargs)
        self.ds1 = Downsample(64, False)
        self.ds2 = Downsample(128)
        self.ds3 = Downsample(256)
        self.ds4 = Downsample(512)
        self.ds5 = Downsample(512)
        self.ds6 = Downsample(512)
        self.ds7 = Downsample(512)
        self.ds8 = Downsample(512)
        self.us1 = Upsample(512, True)
        self.us2 = Upsample(512, True)
        self.us3 = Upsample(512, True)
        self.us4 = Upsample(512)
        self.us5 = Upsample(256)
        self.us6 = Upsample(128)
        self.us7 = Upsample(64)
        self.final = layers.Conv2DTranspose(3, 4, strides=2, padding='same',
                                         kernel_initializer = tf.random_normal_initializer(mean=0., ↴
stddev=0.02),
                                         activation='tanh')
        self.concat = layers.concatenate()

    def call(self, inputs):

```

```

x1 = self.ds1(inputs)
x2 = self.ds2(x1)
x3 = self.ds3(x2)
x4 = self.ds4(x3)
x5 = self.ds5(x4)
x6 = self.ds6(x5)
x7 = self.ds7(x6)
x8 = self.ds8(x7)
x9 = self.us1(x8)
x9 = self.concat([x9, x7])
x10 = self.us2(x9)
x10 = self.concat([x10, x6])
x11 = self.us3(x10)
x11 = self.concat([x11, x5])
x12 = self.us4(x11)
x12 = self.concat([x12, x4])
x13 = self.us5(x12)
x13 = self.concat([x13, x3])
x14 = self.us6(x13)
x14 = self.concat([x14, x2])
x15 = self.us7(x14)
x15 = self.concat([x15, x1])
out = self.final(x15)
return out

```

## 1.8 The Discriminator

Simple Discriminator with 3 downsample units and 2 convolution layers. The output is a  $30 \times 30 \times 1$  patch instead of a single number.

```
[ ]: class Discriminator(keras.models.Model):
    def __init__(self, **kwargs):
        super(Discriminator, self).__init__(**kwargs)
        self.ds1 = Downsample(64, False)
        self.ds2 = Downsample(128)
        self.ds3 = Downsample(256)
        self.zeropad = layers.ZeroPadding2D()
        self.conv = layers.Conv2D(512, 4, strides=1,
                               kernel_initializer = tf.random_normal_initializer(mean=0., □
                               stddev=0.02),
                               use_bias=False)
        self.norm = tfa.layers.InstanceNormalization(
            gamma_initializer = keras.initializers.RandomNormal(mean=0., □
            stddev=0.02))
        self.activation = layers.LeakyReLU()
        self.final = layers.Conv2D(1, 4, strides=1,
```

```

        kernel_initializer = tf.random_normal_initializer(mean=0., stddev=0.02))

    def call(self, inputs):
        x = self.ds1(inputs)
        x = self.ds2(x)
        x = self.ds3(x)
        x = self.zeropad(x)
        x = self.conv(x)
        x = self.norm(x)
        x = self.activation(x)
        x = self.zeropad(x)
        x = self.final(x)
        return x

```

## 1.9 Compelete CycleGAN object

There are 2 generators, one for making real photos look like monet and the other for making monet painting look like real photos.

There are 2 discriminators as well, one for each generator.

There are 4 losses defined here: 1. Descriminator loss: Did the discriminator succesfully determined whether the input was fake or real? 2. Generator loss: Did the generator succesfully fool the discriminator into thinking the image it generated was real? 3. Cycle loss:

real → monet\_gen → fake\_monet → real\_gen → real

and

monet → real\_gen → fake\_real → real\_gen → monet Basically, an input going around the cycle should come out as itself. 4. Identity loss:

monet → monet\_gen → monet

real → real\_gen → real

Overall loss is the sum of these.

```
[ ]: class CycleGAN(keras.models.Model):
    def __init__(self, **kwargs):
        super(CycleGAN, self).__init__(**kwargs)
        self.monet_gen = Generator()
        self.photo_gen = Generator()
        self.monet_dsc = Discriminator()
        self.photo_dsc = Discriminator()
        self.monet_gen_optimizer = keras.optimizers.Adam(2e-4, beta_1=0.5)
        self.photo_gen_optimizer = keras.optimizers.Adam(2e-4, beta_1=0.5)
        self.monet_dsc_optimizer = keras.optimizers.Adam(2e-4, beta_1=0.5)
        self.photo_dsc_optimizer = keras.optimizers.Adam(2e-4, beta_1=0.5)
```

```

def dsc_loss(self, real, generated):
    real_loss = keras.losses.BinaryCrossentropy(from_logits=True,
                                                reduction=keras.losses.Reduction.NONE)(tf.ones_like(real), ↴
    ↪real)
    fake_loss = keras.losses.BinaryCrossentropy(from_logits=True,
                                                reduction=keras.losses.Reduction.NONE)(tf.
    ↪zeros_like(generated), generated)
    return real_loss + fake_loss

def gen_loss(self, generated):
    return keras.losses.BinaryCrossentropy(from_logits=True,
                                            reduction=keras.losses.Reduction.NONE)(tf.ones_like(generated), ↴
    ↪generated)

def cycle_loss(self, real, cycled, lam=10):
    return lam * tf.reduce_mean(tf.abs(real - cycled))

def identity_loss(self, real, same, lam=10):
    return lam * 0.5 * tf.reduce_mean(tf.abs(real - same))

def train_step(self, batch):
    photo, monet = batch
    with tf.GradientTape(persistent=True) as tape:
        fake_monet = self.monet_gen(photo)
        cycled_photo = self.photo_gen(fake_monet)

        fake_photo = self.photo_gen(monet)
        cycled_monet = self.monet_gen(fake_photo)

        same_photo = self.photo_gen(photo)
        same_monet = self.monet_gen(monet)

        disc_real_monet = self.monet_dsc(monet)
        disc_real_photo = self.photo_dsc(photo)

        disc_fake_monet = self.monet_dsc(fake_monet)
        disc_fake_photo = self.photo_dsc(fake_photo)

        monet_dsc_loss = self.dsc_loss(disc_real_monet, disc_fake_monet)
        photo_dsc_loss = self.dsc_loss(disc_real_photo, disc_fake_photo)

        total_cycle_loss = self.cycle_loss(monet, cycled_monet) + self.
    ↪cycle_loss(photo, cycled_photo)

        monet_gen_loss = self.gen_loss(disc_fake_monet) + total_cycle_loss
    ↪+\

        self.identity_loss(monet, same_monet)

```

```

    photo_gen_loss = self.gen_loss(disc_fake_photo) + total_cycle_loss
    ↵+\

        self.identity_loss(photo, same_photo)

    m_gen_grad = tape.gradient(monet_gen_loss, self.monet_gen.
    ↵trainable_variables)
    m_dsc_grad = tape.gradient(monet_dsc_loss, self.monet_dsc.
    ↵trainable_variables)
    p_gen_grad = tape.gradient(photo_gen_loss, self.photo_gen.
    ↵trainable_variables)
    p_dsc_grad = tape.gradient(photo_dsc_loss, self.photo_dsc.
    ↵trainable_variables)

    self.monet_gen_optimizer.apply_gradients(zip(m_gen_grad, self.monet_gen.
    ↵trainable_variables))
    self.monet_dsc_optimizer.apply_gradients(zip(m_dsc_grad, self.monet_dsc.
    ↵trainable_variables))
    self.photo_gen_optimizer.apply_gradients(zip(p_gen_grad, self.photo_gen.
    ↵trainable_variables))
    self.photo_dsc_optimizer.apply_gradients(zip(p_dsc_grad, self.photo_dsc.
    ↵trainable_variables))

    return {
        "m_gen_loss": monet_gen_loss,
        "p_gen_loss": photo_gen_loss,
        "m_dsc_loss": monet_dsc_loss,
        "p_dsc_loss": photo_dsc_loss,
    }

def train(self, ds, epochs):
    for epoch in range(epochs):
        for data in ds:
            losses = self.train_step(data)
            tf.print(f"""epoch: {epoch},\t m_gen loss: {tf.
            ↵reduce_mean(losses["m_gen_loss"])},\t p_gen loss: {tf.
            ↵reduce_mean(losses["p_gen_loss"])},\t m_dsc loss: {tf.
            ↵reduce_mean(losses["m_dsc_loss"])},\t p_dsc loss: {tf.
            ↵reduce_mean(losses["p_dsc_loss"])}""")

    return None

```


Create the object and test an input

```
[ ]: with strategy.scope():
    cycle_gan_model = CycleGAN()
    test_monet = cycle_gan_model.monet_gen(example_photo)
```

```
[ ]: fig, ax = plt.subplots(1, 2, dpi=150)
ax[0].set_title('Photo')
ax[0].imshow(example_photo[0]*0.5 + 0.5)

ax[1].set_title('Monet')
ax[1].imshow(test_monet[0]*0.5 + 0.5)
```

```
[ ]: <matplotlib.image.AxesImage at 0x7f53ec208fa0>
```



Outputs noise because it has random weights before training.

## 1.10 Train

Train CycleGAN for 10 epochs. Notice the two datasets are zipped and passed to the model. It doesn't matter which two pair of photo/monet are passed at each epoch as long as the same ones don't get passed in. Also the batch size for both data sets is 1 so each epoch is just one pair of photos.

```
[ ]: with strategy.scope():
    cycle_gan_model = CycleGAN()
    cycle_gan_model.train(
        tf.data.Dataset.zip((photo_ds, monet_ds)),
        epochs=10
    )
```

```
epoch: 0,          m_gen loss: 4.036203384399414,  p_gen loss: 5.0994553565979,
m_dsc loss: 1.2989656925201416,  p_dsc loss: 0.791059136390686
```

```

epoch: 1,          m_gen loss: 3.5121445655822754,          p_gen loss:
4.297967433929443, m_dsc loss: 1.546463131904602, p_dsc loss: 1.2357348203659058
epoch: 2,          m_gen loss: 3.261470079421997,  p_gen loss: 4.106157302856445,
m_dsc loss: 1.4388259649276733, p_dsc loss: 1.1355243921279907
epoch: 3,          m_gen loss: 3.1116626262664795,          p_gen loss:
3.7555954456329346, m_dsc loss: 1.3309767246246338, p_dsc loss:
1.3365322351455688
epoch: 4,          m_gen loss: 2.8883275985717773,          p_gen loss:
3.7058777809143066, m_dsc loss: 1.2498728036880493, p_dsc loss:
1.317004680633545
epoch: 5,          m_gen loss: 2.8293566703796387,          p_gen loss:
3.718421697616577, m_dsc loss: 1.2094194889068604, p_dsc loss: 1.297163724899292
epoch: 6,          m_gen loss: 2.8371212482452393,          p_gen loss:
3.633399486541748, m_dsc loss: 1.1624222993850708, p_dsc loss: 1.247989535331726
epoch: 7,          m_gen loss: 2.814274311065674,  p_gen loss: 3.5465259552001953,
m_dsc loss: 1.1558595895767212, p_dsc loss: 1.2297011613845825
epoch: 8,          m_gen loss: 2.905796766281128,  p_gen loss: 3.956342935562134,
m_dsc loss: 1.1680665016174316, p_dsc loss: 0.9180918335914612
epoch: 9,          m_gen loss: 2.9855830669403076,          p_gen loss:
3.655225992202759, m_dsc loss: 1.2676938772201538, p_dsc loss:
1.0644094944000244

```

I couldn't afford to train for longer but 10 epochs already shows improvements

```

[ ]: _, ax = plt.subplots(5, 2, figsize=(12, 12), dpi=150)
for i, img in enumerate(photo_ds.take(5)):
    prediction = cycle_gan_model.monet_gen(img, training=False)[0].numpy()
    prediction = (prediction * 127.5 + 127.5).astype(np.uint8)
    img = (img[0] * 127.5 + 127.5).numpy().astype(np.uint8)

    ax[i, 0].imshow(img)
    ax[i, 1].imshow(prediction)
    ax[i, 0].set_title("Input Photo")
    ax[i, 1].set_title("Monet-esque")
    ax[i, 0].axis("off")
    ax[i, 1].axis("off")
plt.show()

```

Input Photo



Monet-esque



Input Photo



Monet-esque




Input Photo




Monet-esque



Input Photo




Monet-esque



Input Photo



Monet-esque



You can see the texture added on top of the pictures here to make them look more like a painting!

More training steps should give better results too!

[ ]: